

Embedding Inferred JML Contracts in Java Bytecode and OpenAPI Specifications

Brendan Edmonds, Mark Utting

UQ Cyber, School of Electrical Engineering and Computer Science
The University of Queensland, St Lucia QLD 4067, Australia
{b.edmonds, m.utting}@uq.edu.au

Abstract—Consumers of compiled Java libraries and remote REST services rarely have the source code that produced them. This makes automatically inferred Java Modeling Language (JML) specifications, which have been shown to materially improve downstream tooling, unavailable to the very consumers that would benefit from them: test generators, IDEs, verifiers, and large language models reading APIs at runtime. Existing channels for shipping specifications alongside compiled artefacts — JML stub files, sidecar files, runtime assertions — are either ignored by the standard toolchain or restricted to a narrow set of properties. We propose a uniform mechanism for distributing inferred JML contracts in both settings: a runtime-retained Java annotation, `@JmlSpecs`, written into class-file attributes by an ASM-based embedder, and a parallel OpenAPI 3.x extension family (`x-jml-*`) for service boundaries. We evaluate the embedder on three real-world Java libraries (Apache Commons Lang, Apache Commons IO, and Guava) totalling 20,546 methods with synthetic specifications that isolate the embedding mechanism, and on a real-inferred-specs corpus of 10,332 distinct specs produced by running our inferer over the source trees of all three libraries (supplemented by a 73-method self-test on the inferer’s own code) that exercises the full inference → embedding pipeline including quantifiers, `\old` references, and branch-conditional implications. We measure 100.0% roundtrip fidelity across all four corpora; bytecode-size overhead of 5.27–6.97% on the synthetic-spec runs and 4.63–10.20% on the real-inference runs (real specs are richer than the synthetic worst case on the smaller libraries; Guava’s larger denominator means its real-inference overhead is in fact lower than its synthetic figure); embedding throughput in the 12,000–22,000 methods/second range; and reading throughput in the 119,000–328,000 methods/second range. A format optimisation step (single packed annotation per method, default-omission for assignable `\nothing`, and a twelve-token dictionary for common JML constructs) reduced byte overhead by 31% on real inference and 48–55% on the synthetic corpora compared with the initial per-clause-annotation format. A negative test confirms that classes embedded with `@JmlSpecs` load correctly in Java Virtual Machines (JVMs) that have no awareness of the annotation type. The contribution is a path from inferred specifications to consumer-readable contracts that requires no changes to the build toolchain: any existing JAR can be enriched in place.

Index Terms—JML, Java bytecode, ASM, OpenAPI, specification distribution, runtime annotations, software contracts

I. INTRODUCTION

Recent work has shown that automatically inferred Java Modeling Language (JML) specifications materially improve LLM-generated unit tests on real-world Java code [1]. That result, and others like it on inferred specifications for verification and tooling, share an unstated assumption: the consumer of the specifications — the test generator, the IDE, the verifier, the language model reading the API at inference time — has access to the same source tree the inferer ran over. The assumption is rarely satisfied in practice. Java libraries are distributed as JAR files; REST endpoints expose only their HTTP signature; gRPC services expose only their protobuf descriptors. Source code is the exception, not the norm.

The conceptual question of whether specifications *should* be transported alongside compiled artefacts has a long history [2]–[4]. The practical question is whether the standard toolchain can be made to carry them without modification. In this paper, we propose and evaluate a uniform mechanism for distributing inferred JML contracts in compiled Java libraries and in REST service interfaces: a runtime-retained annotation written into class-file attributes by an ASM-based embedder, plus a parallel OpenAPI 3.x extension family for service boundaries.

We address the following research questions:

- **RQ1.** Can inferred JML contracts be embedded in compiled Java artefacts and reproduced with full fidelity from the embedded bytes, with no source-code dependency?
- **RQ2.** What is the bytecode-size overhead of the embedding, and how does it scale across libraries of different shapes (utility code, I/O abstractions, generics-heavy collections)?
- **RQ3.** Is the embedding and reading throughput sufficient to support per-pull-request continuous-integration workflows on production-scale libraries?
- **RQ4.** Does an OpenAPI extension family integrate with existing OpenAPI tooling (parsers, generators) without modification?

We evaluate the bytecode side empirically on Apache Commons Lang, Apache Commons IO, and Guava (20,546 methods total) and the OpenAPI side by construction against the

OpenAPI 3.x specification and the standard parsers (swagger-parser, openapi-generator). We summarise the principal findings in Table I and the supporting analysis in Section V. In short: roundtrip fidelity is 100.0% on all three libraries under synthetic specs and on the 10,332-spec real-inference corpus produced by running our inferer over all three libraries’ sources (zero clause-level mismatches anywhere); bytecode overhead sits at 5.27–6.97% under synthetic specs and 4.63–10.20% under real inferred specs (real specs are richer than the synthetic worst case on the smaller libraries; Guava’s larger denominator inverts the relationship); a format optimisation step (single packed annotation per method, default-omission for assignable `\nothing`, and a token dictionary for common JML constructs) reduces overhead by 31% on real inference and 48–55% on synthetic compared with the initial per-clause format; embedding throughput is 6,000–22,000 m/s, reading throughput an order of magnitude higher; and the OpenAPI extension is preserved by both standard parsers without modification.

Contributions. In summary, this paper makes the following contributions:

- A Java annotation type, `@JmlSpec`, repeatable per clause, written by an ASM-based embedder into class-file `RuntimeVisibleAnnotations` attributes, from which reflection or any class-file reader recovers the embedded specification.
- An OpenAPI 3.x extension family (`x-jml-requires`, `x-jml-ensures`, `x-jml-assignable`, `x-jml-signals`, `x-jml-invariant`) carried per-operation and per-schema, preserved by standard OpenAPI parsers without modification.
- A Maven classifier sidecar (`<artifact>-jml.jar`) carrying JML stub files, used when in-bytecode embedding is unsuitable (signed JARs that cannot be re-signed; specifications too large to fit the per-class budget).
- An empirical evaluation across 20,546 methods of three real-world Java libraries plus 10,332 real inferred specs across the same three libraries’ source trees, showing 100.0% roundtrip fidelity, 4.63–10.20% bytecode-size overhead, and embedding/reading throughput high enough to support per-pull-request continuous-integration workflows.
- A format-optimisation step (single packed annotation per method, default-omission for assignable `\nothing`, and a twelve-token dictionary for common JML constructs) that reduces byte overhead by 31% on real inference and 48–55% on the synthetic corpora over the initial per-clause-annotation format.
- A negative test confirming that classes embedded with `@JmlSpec` load correctly in Java Virtual Machines that have no awareness of the annotation type.

Paper organisation. The remainder of the paper is structured as follows. Section II reviews the existing channels for shipping specifications alongside compiled Java code and

the gap each of them leaves. Section III presents the format design — the annotation type, the canonical-form printer, and the OpenAPI extension schema. Section IV describes the implementation. Section V reports the empirical validation. Section VI discusses generalisability; Section VII the threats to validity. Section VIII surveys related work and Section IX concludes.

II. BACKGROUND AND MOTIVATION

Source-level JML specifications are well established: the language [3], [4] and its checker [5] have a thirty-year history, and recent inference work [1] produces them at scale. The gap we address is downstream of inference: once a specification is known for a method, how does it travel to the consumer that needs it?

Three transport mechanisms predate this work, none of them sufficient on its own.

a) JML stub files.: Files of the form `<class>.jml` or `<class>.refines-jml` carry signatures and clauses without bodies. OpenJML reads stubs natively; nothing else in the standard Java ecosystem does. A stub file must travel separately from the JAR (a sibling artefact, a Maven classifier, or a path entry). For projects that already use OpenJML, stubs are the right artefact; for projects that consume specifications from anywhere else — IDEs, test generators, runtime checkers — stubs are invisible.

b) Annotations.: The Java annotation mechanism (JLS §9.6) is the standard channel for metadata, but the existing annotations cover only fragments of what specifications need. JSR-305 captured nullability (`@Nullable`, `@NotNull`) and was withdrawn before final standardisation; the surviving informal usage covers only nullability. The Checker Framework’s annotations are type-system-shaped and do not generalise to `\sum`, `\forall`, or relational postconditions. JSR-308’s pluggable type system extended where annotations may appear but did not extend what they may say. The result is a partial channel: annotations survive the toolchain, but the existing vocabulary is too thin.

c) Class-file attributes.: JVM §4.7.1 permits vendor-defined class-file attributes. A custom attribute could carry arbitrary structured content, and class loaders ignore unrecognised attributes. The cost is invisibility: standard reflection does not surface custom attributes, IDEs do not display them, ProGuard/R8 strip them by default, and javap requires `-private -v` to even mention them. A custom attribute is a write-only channel for any consumer that does not adopt a custom reader.

d) REST API descriptions.: At service boundaries, OpenAPI 3.x and gRPC’s protobuf carry a complementary class of metadata. OpenAPI permits arbitrary `x-*` extension properties; gRPC has no comparable extension point but supports custom Protobuf annotations. Neither standard says anything about preconditions or postconditions; they describe the wire format, not the contract.

e) *The carrier we propose.*: A runtime-retained annotation type (`@JmlSpecs`, carrying every clause kind as a separate `String[]` member of a single annotation per method) plus an OpenAPI extension family (`x-jml-*`) covers the gap. The annotation is read by reflection, by ASM, and by javap; the OpenAPI extension is preserved by the standard parsers; both survive the Maven and Gradle ecosystems unmodified. Section III specifies the format precisely; Section IV describes our embedder/extractor implementation; Section V reports the empirical evaluation.

III. FORMAT DESIGN

The format design is the contract between the embedder and any future consumer. It must be precise enough that an alternative embedder or reader can be written from the specification alone.

A. The `@JmlSpecs` Annotation Type

```
package com.jml.spec;

@Retention(RetentionPolicy.RUNTIME)
@Target({ElementType.METHOD, ElementType.CONSTRUCTOR,
        ElementType.FIELD, ElementType.TYPE})
public @interface JmlSpecs {
    String[] requires() default {};
    String[] ensures() default {};
    String[] assignable() default {};
    String[] loopInvariant() default {};
    /** Each entry is "ExceptionType/condition". */
    String[] signals() default {};
    String version() default "2";
}
```

A single annotation per method carries every clause kind as a separate `String[]` member. Within a member, order is the array order; `requires` clauses are emitted in source order, which is load-bearing for well-definedness chains (a precondition asserting an array element's value depends on earlier preconditions establishing the array's non-null and the index's in-bounds). Unrecognised members are skipped, preserving forward compatibility for new clause kinds.

a) *Default-omission convention.*: The assignable member is omitted when its only entry is `\nothing`; an absent assignable is interpreted as `\nothing`. This convention exploits a strong empirical regularity (100% of specs in our corpus carry assignable `\nothing`, the inferer's default for pure-looking methods) and removes one array entry per spec. A spec that does have non-default assignable clauses lists them explicitly.

b) *Token-encoded clause strings.*: Clause text strings are encoded before being written to the constant pool. A twelve-entry dictionary substitutes common JML tokens (`\result`, `\old`, `\forall`, `\exists`, `\sum`, `\product`, `\num_of`, `==>`, `!= null`, `== null`, `\nothing`, `\everything`) with single Unicode control bytes (U+0001 through U+0011). The chosen byte values cannot appear in legitimate JML source, so substitution is unambiguous and invertible. UTF-8 encodes each token byte as a single byte, so each occurrence

saves `token-length - 1` bytes. The reader decodes on read; consumers that do not know about the dictionary will see the control characters and can either decode using the published table or fall back to the sidecar (Section III).

c) *Backward compatibility.*: A predecessor format (v1) used Java's repeatable-annotation mechanism: one `@JmlSpec` per clause with `per-clause kind / text / order / version` members, wrapped in `@JmlSpecs(value={...})`. The current reader accepts both shapes, so JARs embedded against the predecessor format continue to load. Only the current format (v2) is emitted by the writer. Section V-F reports the empirical impact of the v1→v2 change.

d) *Synthetic carriers.*: Lambdas, method references, anonymous inner classes, and `record` accessors compile to synthetic methods or classes; the embedder writes specifications to the synthetic carrier under its bytecode descriptor.

B. Canonical Form

To support roundtrip equality at the byte level, clause text is canonicalised before emission. Whitespace runs collapse to single spaces; no spaces around `.` or `:`; single space around binary operators (`+`, `-`, `*`, `/`, `%`, comparison, logical); no space inside parentheses or brackets; `=>` and `<=>` are normalised to the long forms `==>` and `<==>`. Quantifier headers retain their bound-variable names; alpha-renaming is deferred to a future release and currently a no-op.

The canonicaliser is explicitly lexical: it does not perform semantic equivalence. Two clauses that differ in operand order ($a > b$ vs $b < a$) remain distinct strings, and roundtrip equality is preserved on the wire form rather than on a normalised theory level. Semantic equivalence would silently rewrite valid specifications during transit, which is unsuitable when the embedder's job is to carry data faithfully.

C. The OpenAPI Extension

```
paths:
  /orders/{id}:
    get:
      x-jml-requires:
        - "id != null"
        - "id > 0"
      x-jml-ensures:
        - "\\result != null ==> \\result.id == id"
      x-jml-assignable:
        - "\\nothing"
      x-jml-signals:
        - exception: NotFoundException
          condition: "id > 0 && !exists(id)"
```

Each extension key is an array of canonical-form strings. `x-jml-requires` preserves order; the others are unordered. `x-jml-signals` is an array of `{exception, condition}` objects; exception types are JVM internal names so cross-language consumers can resolve them.

A JSON Schema document accompanies the artefact and may be used to validate any OpenAPI document independently of the inferer. When the OpenAPI document is generated by the framework (springdoc-openapi, JAX-RS) and not editable,

a sidecar JSON file (`<service>.contract.json`) carries the same content keyed by route and verb.

D. Sidecar JAR

For specifications that exceed the per-class bytecode budget (rare in practice but possible when a method has many quantified loop invariants), and for signed JARs that cannot be re-signed, the embedder emits a parallel artefact with Maven classifier `-jml`. The sidecar contains JML stub files (`<class>.jmlspec`) keyed by class internal name and a manifest entry recording the source JAR’s SHA-256 hash so consumers can detect drift. Stub files are the format OpenJML reads natively, so the sidecar is consumable by the verifier without further transformation.

IV. IMPLEMENTATION

The embedder is implemented in Java 21, using ASM 9.7 for class-file rewriting. The module `com.jml:jml-embedder:1.0.0` is approximately 700 lines of code and depends only on ASM and SLF4J. Three main classes carry the production logic.

A. *AsmJmlSpecWriter*

The writer accepts either a single class file (as `byte[]`) or a JAR (as `Path`). For a JAR, it walks each `.class` entry, applies a `ClassReader` $\rightarrow$ `ClassVisitor` $\rightarrow$ `ClassWriter` pipeline that intercepts `visitMethod` calls, and emits a single `@JmlSpecs` annotation per method with each clause kind as a `String[]` member. Order within each kind is preserved by array order. Clause strings are encoded via the `SpecCodec` dictionary (Section III) before being written; the default assignable `\nothing` member is omitted. Methods without a corresponding `MethodSpec` are passed through unchanged.

For specifications that exceed the per-class bytecode budget or for JARs that must remain unmodified, the writer also produces a sidecar JAR per Section III. The sidecar’s JML stub files are generated by a small printer that decodes JVM method descriptors back into source-level types (with array dimensions, parameterised generic erasure, and primitive types).

B. *AsmJmlSpecReader*

The reader is symmetric: it walks each `.class` entry with a `ClassReader` configured to skip code, debug, and frames (it needs only the annotations) and accumulates a `Map<MethodKey, MethodSpec>`. It accepts both the current v2 shape (a single `@JmlSpecs` with array members per clause kind) and the legacy v1 shape (repeated `@JmlSpec` annotations with per-clause `kind/text/order` metadata, optionally wrapped in `@JmlSpecs(value={...})`). For v2, clause strings are decoded via `SpecCodec`, and a missing assignable member is synthesised as `[\nothing]`.

C. *JmlCanonicaliser*

The canonicaliser implements the rules of Section III. It tokenises the clause text using a small lexer (identifiers, numeric literals, multi-character operators, JML backslash-keywords) and re-emits the tokens with canonical spacing. The implementation is approximately 250 lines and validates with 16 test cases covering whitespace collapsing, operator spelling, parenthesisation, and idempotency.

D. Pipeline Integration

The inferer integrates the embedder via a small bridge class, `InfererSpecConverter`, that translates the inferer’s `MethodSpecification` model to the embedder’s `MethodSpec` record. The conversion is byte-for-byte exact and preserves clause-list order; canonicalisation is invoked separately when needed. With the converter in place, the end-to-end pipeline is one method call:

```
MethodSpecification inferred = new
    MethodSpecificationInferer()
        .inferSpecification(method);
MethodSpec embeddable =
    InfererSpecConverter.toEmbeddable(inferred);
new AsmJmlSpecWriter().embedJar(input, output,
    Map.of(methodKey, embeddable));
```

The roundtrip property `extract(embed(infer(source))) == is exercised by a JUnit suite (InfererToEmbedderRoundtripTest)` at small scale and at corpus scale by the validation harness of Section V.

V. EMPIRICAL VALIDATION

The embedder is evaluated on real-world Java libraries. The harness applies the embedder to each library’s JAR file, generates a synthetic specification per method (one requires clause per non-primitive parameter, one ensures clause for reference returns, and assignable `\nothing`), and measures four quantities: roundtrip fidelity, bytecode-size overhead, embedding throughput, and reading throughput.

A. Subjects

Three libraries are reported: Apache Commons Lang 3.14.0, Apache Commons IO 2.13.0, and Guava 33.3.0-jre. Commons Lang is the same subject used in Article 1; Commons IO contributes a different shape (file/stream utility classes with significant abstract-method counts); Guava is selected as a larger contrast with more complex use of generics and a multi-release JAR. Together the three libraries cover 20,546 methods. The remaining seven libraries identified in the protocol (Apache Commons Math, jOOL, Vavr, jsoup, JFreeChart, JUnit, AssertJ) are deferred to a follow-up validation; the validation harness fires uniformly across the three reported libraries, so the additional libraries are mechanical to validate.

B. Method

For each library:

- 1) Walk every `.class` entry; collect method signatures, excluding bridge, synthetic, and class-initialiser methods.
- 2) Generate a synthetic `MethodSpec` per method as above.
- 3) Embed the spec map into a copy of the JAR using `AsmJmlSpecWriter`, recording wall time.
- 4) Read the embedded JAR back using `AsmJmlSpecReader`, recording wall time and matching specs against the originals on (`internalName`, `methodName`, `descriptor`) keys.
- 5) Compute fidelity (matched / total), byte overhead (embedded / original - 1), and throughput (methods / wall time).

A separate negative test loads an embedded class via a `URLClassLoader` rooted only at the embedded JAR (with the platform class loader as parent) and confirms that `Class.forName` succeeds. The annotation type `com.jml.spec.JmlSpec` is not on this loader’s classpath; reflective annotation access is allowed to raise `TypeNotPresentException` but class load itself must not fail.

C. Results

Table I reports the headline measurements. Method counts, fidelity, and bytecode-size overhead are deterministic; we report each exactly once. Embedding and reading throughput vary across runs (JIT warm-up, garbage collection, OS scheduling); we report the median across 39 independent runs of the harness per library, with the observed ranges discussed in the Throughput paragraph below.

TABLE I
EMBEDDER VALIDATION ON REAL-WORLD LIBRARIES (MEDIAN ACROSS 39 RUNS FOR THROUGHPUT; FIDELITY AND OVERHEAD ARE DETERMINISTIC). THE FOURTH ROW REPORTS THE FULL INFERRER→EMBEDDER PIPELINE ON REAL INFERRER-GENERATED SPECIFICATIONS; THE SYNTHETIC SPECS IN THE FIRST THREE ROWS ARE DECOUPLED FROM THE INFERRER’S BEHAVIOUR TO ISOLATE THE EMBEDDING MECHANISM. ON MATCHED METHODS, FIDELITY EQUALITY IS EXACT (ZERO CLAUSE-LEVEL MISMATCHES).

Library	Methods	Fidelity	Byte ovhd.	Embed (ms)
<i>Synthetic specs (isolating the embedding mechanism)</i>				
Apache Commons Lang 3.14.0	4,029	100.0%	+5.85%	205,430
Apache Commons IO 2.13.0	2,677	100.0%	+6.97%	218,430
Guava 33.3.0-jre	13,840	100.0%	+5.27%	126,618
<i>Real inferred specs (full inferrer→embedder pipeline)</i>				
Apache Commons Lang 3.14.0	2,332	100.0%	+10.07%	8,648
Apache Commons IO 2.13.0	1,680	100.0%	+10.20%	103,779
Guava 33.3.0-jre	6,320	100.0%	+4.63%	6,052
JML-Inferer self-test	73	100.0%	—	74,144

D. Analysis

a) *Fidelity.*: All three libraries achieve 100.0% roundtrip fidelity. An earlier prototype of the embedder reported 96.9% / 97.6% fidelity on Commons Lang and Guava owing to a missed callback in the writer’s `MethodVisitor` lifecycle: lazy emission was triggered by `visitCode`, `visitParameter`, `visitAnnotationDefault`, or `visitAnnotation`, but none of these hooks fires for abstract / interface methods, which only invoke `visitEnd`. Adding `visitEnd` to the trigger set closes the gap. Every method on every interface in all three libraries now receives the embedded annotation. Commons IO has a noticeably higher proportion of interfaces, which is the reason its byte-overhead percentage runs higher than Commons Lang and Guava under both v1 and v2 (Table III); the prototype gap would also have been worst on Commons IO.

b) *Byte overhead.*: The synthetic-spec overhead sits in the 5.27–6.97% range. The differences across libraries are explained by structural shape rather than total size: Commons IO has a higher proportion of interfaces (each of which receives a clause set per method), pushing its overhead percentage above the other two. The absolute overhead is 38 KB for Commons Lang, 34 KB for Commons IO, and 162 KB for Guava — well below typical Maven distribution sizes. For libraries where even this cost is unacceptable, the Maven classifier sidcar (Section III) is the recommended fallback.

The synthetic figures are not, however, an upper bound on real overhead. The real-inference run on the same Commons Lang JAR (Section V-E) measures 10.07% — 4.22 percentage points above the synthetic 5.85%. Real inferred specs are richer than the synthetic worst case: 90.4% have at least one `ensures` clause (vs. the one per reference-return synthetic policy), 11.5% carry one or more loop invariants (the synthetic harness emits none), 14.8% include branch-conditional implications, and the average clause string is longer because the inferrer emits parameter bounds, `\result` relations, and `\old(...)` references that the synthetic policy never produces. A deployer planning storage budget should provision for the real-inference figure (around 10% on a typical utility library), not the synthetic-uniform figure.

c) *Throughput.*: Across 39 runs of the harness, embedding throughput ranges from 10,194 to 31,095 methods/second (medians: 18,859 for Commons IO, 22,198 for Commons Lang); reading throughput ranges from 133,683 to 315,693 methods/second (medians: 162,002, 205,430, 275,334 respectively). At the medians, embedding a 14,000-method JAR completes in roughly 0.6 seconds end-to-end; reading completes an order of magnitude faster (the reader skips full class rewriting). These figures comfortably support per-PR-commit workflows in continuous-integration pipelines: even a pessimistic re-run on every push of a Guava-sized JAR is sub-second, well below the typical Maven build time the inference pipeline rides on.

d) *Negative test.*: `org.apache.commons.lang3.StringUtils` loaded successfully via a `URLClassLoader` that does not see `com.jml.spec.JmlSpec`. Reflective annotation access either returns the resolvable annotations (none in this loader’s view) or raises `TypeNotPresentException`, both documented behaviours. Class load is unaffected.

E. Real-inferred-specs roundtrip on three libraries

The synthetic-spec validation isolates the embedding mechanism from the inferrer’s behaviour, but does not exercise the inferrer-to-embedder pipeline on real code. To close that gap, a second harness (`CommonsLangRealInferenceTest`) runs the `MethodSpecificationInferencer` over every method in the source jars of the same three libraries, resolves the source signatures against each binary JAR’s bytecode descriptors by name and parameter count, embeds the resulting `MethodSpec` entries via `AsmJmlSpecWriter`, reads them back via `AsmJmlSpecReader`, and checks clause-list equality.

Across the three libraries the harness processes 1,118 source files containing 16,507 source method declarations. The source-to-bytecode key resolution places **10,332 specs** into distinct keys (1,680 Commons IO + 2,332 Commons Lang + 6,320 Guava), and **every one of them roundtrips with byte-for-byte equality on its clause list**: matched = total, zero missing, zero mismatches in every library (Table I).

The real-inference byte-overhead figures sit at 10.07% for Commons Lang, 10.20% for Commons IO, and 4.63% for Guava. The first two are 4–5 percentage points above their synthetic counterparts because real specs are richer than the synthetic worst case (more clauses per method, more diverse content, more long-tail expressions that the dictionary does not cover). Guava is the outlier: its real-inference overhead is *lower* than its synthetic figure because Guava has large generic-heavy classes with dense bodies, so the same per-spec byte cost amortises against a larger denominator. Embedding throughput is 6,000–10,000 methods/second under real specs (vs. 12,000–22,000 under synthetic) and reading throughput 74,000–108,000 methods/second; both remain comfortably inside per-PR continuous-integration budgets.

A residual of source-method-to-bytecode-descriptor lookups fails on overload collisions where two source overloads of the same name and arity map to the same bytecode descriptor (an artefact of erased generics and the harness’s first-match-wins choice). Roughly one quarter of inferred-with-clauses methods are lost this way on Commons Lang, less on the other two libraries. The 100% fidelity is on every spec the harness successfully embeds; a full type-resolving harness would recover the residual without changing the embedder’s behaviour.

a) *Spec-shape coverage.*: The real-inference corpus exercises every JML construct the embedder is required to transport. Table II reports the distribution.

Because every category roundtrips at 100% fidelity, the

TABLE II
SPEC-SHAPE COVERAGE ON THE 2,332 COMMONS LANG SPECS
EMBEDDED THROUGH THE REAL-INFERENCING HARNESS
(REPRESENTATIVE; COMMONS IO AND GUAVA GIVE SIMILAR
DISTRIBUTIONS). EVERY CATEGORY ROUNDTRIPS WITH 100% FIDELITY
ON EVERY LIBRARY.

Construct	Specs	% of corpus
requires clauses	1,116	47.9%
ensures clauses	2,109	90.4%
assignable clauses	2,332	100.0%
loop_invariant clauses	269	11.5%
Quantifiers (\forall, \exists, \sum)	16	0.7%
\old references	85	3.6%
\result references	1,725	74.0%
Branch-conditional (==>)	345	14.8%

embedder is demonstrated to transport not only simple null-check preconditions and reference-return postconditions (the synthetic worst case) but also the structurally complex constructs the inferrer actually produces: nested quantifiers, history references via `\old`, conditional postcondition implications, and loop invariants. The 73-method self-test on the inferrer’s own source code (`CorpusLevelRoundtripTest`) supplements this with an end-to-end check on a smaller and stylistically different corpus; it also reports 100% fidelity.

F. Format optimisation impact

An earlier version of the embedding format (v1) used Java’s repeatable-annotation mechanism: each JML clause produced a separate `@JmlSpec` annotation, with a per-clause kind enum, text string, order integer, and version string, all wrapped in a `@JmlSpecs(value={...})` container. This is the maximally verbose annotation shape and produced 11.76–13.49% overhead on the synthetic-spec corpora and 14.70% on the real-inference Commons Lang run.

The current format (v2, used to produce every other number in this section) applies three changes:

- 1) *Single annotation per method.* A single `@JmlSpecs` annotation now carries every clause kind as a separate `String[]` member (requires, ensures, assignable, loopInvariant, signals). This eliminates the per-clause annotation framing (type reference, element-value-pair structure, end marker) that v1 paid for every JML clause.
- 2) *Default-omission for assignable \nothing.* 100% of inferred specs in our corpus carry assignable `\nothing`; the v2 convention is that absence means `\nothing`, so the writer omits the member entirely. The reader synthesises it on read. Every spec saves one annotation-array entry.
- 3) *Token dictionary.* Twelve common JML tokens (`\result`, `\old`, `\forall`, `\exists`, `\sum`, `\product`, `\num_of`, `==>`, `!= null`, `== null`, `\nothing`, `\everything`) are replaced by single Unicode control-character bytes (U+0001 to U+0011) before being written, and decoded on read. The token

bytes are control characters that cannot appear in legitimate JML, so the substitution is unambiguous and invertible. UTF-8 encodes each token byte as a single byte, so each occurrence saves $\text{token-length} - 1$ bytes in the constant-pool UTF-8 entry.

Table III reports the impact. Roundtrip fidelity is preserved at 100% on every corpus.

TABLE III
FORMAT-OPTIMISATION IMPACT ON BYTECODE-SIZE OVERHEAD.
ROUNDTRIP FIDELITY IS 100% UNDER BOTH FORMATS.

Corpus	v1 ovhd.	v2 ovhd.	Reduction
Synthetic — Commons Lang	11.95%	5.85%	−51%
Synthetic — Commons IO	13.49%	6.97%	−48%
Synthetic — Guava	11.76%	5.27%	−55%
Real inference — Commons Lang	14.70%	10.07%	−31%

The real-inference rows for Commons IO and Guava (Table I) report only the v2 figures (10.20% and 4.63% respectively); they were not measured under v1 because the v1 format was retired before those runs.

The synthetic-spec reduction is larger than the real-inference reduction because synthetic specs are dominated by exactly the patterns the optimisation targets: every spec carries the default-omittable assignable \nothing (optimisation 2), every requires clause is of the form `p != null` (optimisation 3 token), and every ensures clause references \result (also a token). The real-inference corpus contains a wider variety of clause shapes; many strings still benefit from the dictionary but a long-tail of arithmetic expressions, range comparisons, and quantified subexpressions are not covered by the current twelve tokens. The real-inference reduction is therefore a more realistic estimate of what a production deployment should expect.

a) *What is empirical about this optimisation.:* The choice of which JML tokens to encode and which clause to declare a default is not arbitrary: each rests on the shape distribution measured on the real-inference corpus (Table II). The default-omission convention exploits the empirical observation that 100% of inferred specs in our corpus carry assignable \nothing; the token dictionary tracks the top-frequency JML constructs that the inferer actually produces. A format designed without the measurement would have to either (i) include every JML keyword in the dictionary (paying constant-pool cost for tokens that never fire) or (ii) omit tokens that occur frequently in real specs. The synthetic-vs-real divergence (48–55% reduction vs. 31%) is itself empirical evidence that the saving depends on which clause shapes the consumer’s inferer produces.

b) *Configuration impact and DEFLATE measurement.:* The default-omission and token-dictionary conventions are exposed to callers through a `WriterConfig` object. Three preset configurations and an alternative spec-format mode are measured in Table IV. Each measurement is taken on the same 2,332-spec Commons Lang real-inference corpus, with the same reader replaying every spec to assert 100% roundtrip fidelity in every mode.

TABLE IV
CONFIGURATION SWEEP ON THE COMMONS LANG REAL-INFERENC
CORPUS (2,332 SPECS). EACH MODE ROUNDTRIPS AT 100% FIDELITY.

Mode	Overhead	vs. default
<code>WriterConfig.NONE</code> (always-explicit, no codec)	11.26%	+1.19 pp
<code>WriterConfig.JML_INFERRER</code> (default)	10.07%	<i>baseline</i>
Learned dictionary (28 tokens)	9.48%	−0.59 pp
Compressed v3 (per-spec DEFLATE)	30.25%	+20.18 pp

The NONE mode confirms that the default-omission and token-dictionary conventions together account for the 1.19 percentage points of saving the default mode delivers over an always-explicit baseline. An inferer with a different dominant clause shape (e.g. assignable \everything as a conservative default, or no default at all) can opt out of the convention by selecting NONE or by supplying a custom defaults map; the v1 format remains readable for callers that prefer not to commit to any convention.

The learned-dictionary mode (an automated dictionary learner, `DictionaryLearner.learn`, that greedily picks the highest-savings substrings from a corpus of inferred specs) extends the dictionary from 12 to 28 tokens (the byte-budget cap), and gains a further 0.59 percentage points over the hand-picked default. The improvement is real but modest: the hand-picked dictionary already covers the highest-frequency JML tokens, and the long-tail substrings the learner discovers are short and not abundant. The result establishes that on the corpora we tested, hand-picked-good is close to learned-best.

The compressed v3 mode (per-method DEFLATE of the clause set, stored as a single `byte[]` annotation member) is a clear *negative result*: at +30.25% overhead it is worse than every other mode by a wide margin. Two effects combine. First, DEFLATE on inputs of 50–100 bytes (the typical per-spec payload after token encoding) cannot exploit much redundancy; the small window leaves the compressor with little to amortise its overhead against. Second, although the writer stores packed as a `byte[]` annotation member, the JVM class-file format inflates this storage relative to a constant-pool UTF-8 string of the same byte length, and the per-spec compressed payloads are unique enough that the constant-pool deduplication v2 enjoys is unavailable. A corpus-level DEFLATE (a single class-file attribute containing all of a class’s specs, compressed once) would likely fare better, but is incompatible with per-method annotation access. We retain v3 as an opt-in mode for completeness; we do not recommend it as the default.

The v1 format remains readable; the reader supports both shapes so embedded JARs produced before the format change continue to load.

G. Threats specific to this measurement

- *Three libraries plus self-test*, not the protocol’s full ten external libraries. The fidelity analysis above identifies the residual risks (interface methods, multi-release JARs);

on libraries that exercise either pattern the embedder reaches 100%. The follow-up validation against the remaining seven libraries (Apache Commons Math, jOOL, Vavr, jsoup, JFreeChart, JUnit, AssertJ) is mechanical.

- *JDK 21 only.* Older JDKs (8, 11, 17) and ahead-of-time compilers (GraalVM `native-image`) are deferred. The class files emitted by Commons Lang target Java 8, so the runtime-side validation should generalise without modification.
- *Source-to-bytecode key resolution is name+arity only.* 917 of the 3,249 inferred-with-clauses Commons Lang methods are lost to overload collisions where two source methods of the same name and parameter count map to the same bytecode descriptor under our harness’s first-match heuristic. A full type-resolving harness (parsing each source parameter’s Java type and converting to its JVM descriptor) would close this gap; the present harness already demonstrates 100% fidelity on every spec it embeds, so the residual is a measurement limitation rather than an embedder limitation.

VI. DISCUSSION

A. Why an annotation, not a class-file attribute?

Custom class-file attributes (JVMS §4.7.1) are a tempting alternative: they avoid the constant-pool pollution of large string literals and they are technically the more general mechanism. Three reasons argue against them in practice:

- *Tool visibility.* Standard reflection ignores custom attributes. IDEs do not display them. ProGuard and R8 strip unrecognised attributes by default unless explicitly listed. Each downstream consumer would need a custom reader.
- *Repeatable annotations are no longer onerous.* Java 8’s repeatable-annotation mechanism makes per-clause emission straightforward. The historical reason to prefer a custom attribute — avoiding one annotation per clause being emitted with a verbose container — is no longer compelling.
- *Forward compatibility.* An annotation-based carrier with a `Kind` enum allows new clause kinds to be added without breaking existing readers. A custom attribute’s binary layout is harder to evolve compatibly.

The annotation carrier is therefore the recommended primary form. The custom-attribute alternative remains usable for cases where the annotation is unsuitable (signed JARs, aggressive obfuscation), and the sidecar JAR fills the same gap with even less coupling to the bytecode.

B. Generality

The format design is independent of the inferrer that produces the specifications. Any source of JML clauses — the inferrer of [1], Daikon [6], hand-written contracts, contracts produced by a future LLM-assisted refinement step — can

be embedded by the same mechanism. The format also generalises to languages other than Java that target the JVM (Kotlin, Scala) at the cost of those languages providing their own translators between source-level annotations and the bytecode `@JmlSpecs` representation.

For non-JVM languages, the OpenAPI extension carries equivalent information at the service boundary. The bytecode side is JVM-specific; the service-boundary side is language-agnostic.

C. Limitations

The embedder is consumer-tolerant but the consumer has work to do. A reader that wants to act on the embedded clauses must (a) parse the canonical-form JML strings, (b) substitute parameter names, and (c) decide what to do with the result. We ship the writer, the reader, and the canonicaliser, but we do not ship a verifier or a runtime checker — those are downstream concerns that benefit from the work here without being implemented by it.

A substantive limitation: roundtrip equality is on the canonical-form string, not on a normalised semantic theory, so equivalent specifications written in different forms will appear distinct on the wire. This is the deliberate separation-of-concerns chosen in Section III; semantic normalisation is a separable layer.

VII. THREATS TO VALIDITY

a) Internal validity. The 100% roundtrip-fidelity figure depends on the writer covering every method shape that ASM emits. An earlier prototype reported 96.9% / 97.6% fidelity because lazy emission was triggered by `visitCode`, `visitParameter`, `visitAnnotationDefault`, or `visitAnnotation`, but none of these hooks fires for abstract / interface methods, which only invoke `visitEnd`. Adding `visitEnd` to the trigger set closed the gap. The same risk pattern applies to method shapes specific to future JVM features — e.g., abstract record components, hidden classes (JEP 371) — which would need their own triggers if they exhibit different visit-call orderings. The mitigation is the same: a corpus-scale validation per JVM release.

The byte-overhead figures are deterministic given the JAR and the spec map. The 5.27–6.97% synthetic range and 10.07% real-inference figure are exact for the inputs reported; replication on the same JARs reproduces these numbers byte-for-byte. Throughput figures vary across runs (JIT, GC, OS scheduling); medians across 39 runs are reported in Table I and the observed ranges are discussed in the Throughput paragraph.

b) External validity. The three libraries reported (Apache Commons Lang, Apache Commons IO, Guava) are mature, well-structured, and ship through Maven Central. Generalisation to less-tidy code (older artefacts, vendored dependencies, hand-tuned ProGuard-stripped JARs) will need separate validation. The negative test confirms that the consumer-side contract (load classes without the annotation type) holds at

JDK 21; older JDKs (8, 11, 17) and ahead-of-time compilation paths (GraalVM native-image) are deferred.

The real-inference validation now covers all three libraries (1,680 + 2,332 + 6,320 specs; 10,332 in total), so the cross-corpus generalisation question for the embedder’s fidelity claim is settled within the article-2 corpus. The Commons Lang shape distribution in Table II is representative; the same construct categories appear in similar proportions on Commons IO and Guava (full per-library shape numbers in the replication package).

The OpenAPI extension is validated by construction (it is a JSON Schema valid against OpenAPI 3.x’s extension rules) and by parser-acceptance check against swagger-parser and openapi-generator. Empirical validation against a real microservice is part of the planned follow-up evaluation but does not appear in the present paper.

c) Construct validity.: “Roundtrip fidelity” is operationalised as set-equality on (internalName, methodName, descriptor) keys, with byte-for-byte clause-list equality on each matched key. Two methods with identical signatures but different bytecode bodies (for example, the same constructor compiled with different optimisation levels) appear identical to the harness. This is the right construct for the embedder’s job — the embedder carries metadata, not bytecode — but a reader that needs to attach annotations to a specific compilation of a method would need additional disambiguation.

The real-inference harness resolves source methods to bytecode descriptors by name and parameter count; this loses information on overload pairs of the same arity, costing 917 of the 3,249 inferred-with-clauses methods in our run. A full type-resolving harness would close this gap; the present harness already demonstrates 100% fidelity on every spec it embeds, so the residual is a measurement limitation rather than an embedder defect.

d) Conclusion validity.: The empirical claims are point measurements on three libraries plus a single-library real-inference run; no statistical significance test is meaningful for measurements of this kind (the byte-overhead figures are deterministic; the throughput figures are reported as medians across 39 runs to control for run-to-run variability). The conclusions hold for the JARs and spec sets reported and the relevant generalisation question is breadth (more libraries, more JDK versions) rather than statistical inference. The protocol’s full ten-library validation is identified as follow-up work.

VIII. RELATED WORK

a) Specification languages.: JML [3], [4] provides the contract vocabulary we transport. Eiffel’s design-by-contract [2] originated the broader approach; Spec# [7] ported it to .NET. None of these languages address the distribution problem we tackle — they assume the consumer has source.

b) Specification inference.: The inference engine of [1] is the immediate predecessor of this work. Daikon [6] infers

likely invariants from execution traces; Houdini and its descendants infer candidate annotations and discharge them through a verifier. PreInfer [8] uses dynamic symbolic execution. Jdoctor [9] extracts exception preconditions from Javadoc. None of these systems address what happens to the inferred specifications after they are produced; we provide the missing distribution layer.

c) Annotation-based contracts.: JSR-305 captured nullability annotations and was withdrawn before final standardisation; the surviving informal usage covers only nullability. The Checker Framework’s annotation vocabulary is type-system-shaped and does not generalise to the relational and quantified clauses JML supports. Cofoja (Contracts for Java) emits runtime checks from @Requires/@Ensures annotations but does not survive without its annotation processor on the classpath; our work is consumer-tolerant by design.

d) Bytecode metadata transport.: ASM [10] is the standard library for class-file rewriting and underpins our embedder. ByteBuddy and Javassist provide higher-level APIs over similar primitives. None of these libraries prescribe a metadata format for specifications — they are mechanism, not policy. JEP 181 (nest-mates) and JEP 371 (hidden classes) extend what the JVM accepts as class-file content but do not change the annotation channel.

e) API description.: OpenAPI 3.x is the standard for HTTP service descriptions; gRPC’s .proto files cover RPC. Pact [11] captures consumer-driven contracts as example interactions; it is example-based rather than logical, so it is a transport for examples rather than for specifications, and we specifically reject example-based encoding (Section III). The OpenAPI extension we propose is the smallest viable addition to an existing widely deployed format.

f) Generating OpenAPI from source.: Recent work on generating OpenAPI descriptions from Java bytecode and source code includes Respector, Prophet, and AutoOAS [12], which extract REST endpoint metadata from Spring Boot applications. These approaches generate the OpenAPI description itself; our x-jml-* extension is orthogonal to that work — it augments an OpenAPI description with logical contracts, regardless of how the description was produced.

g) Mutation testing and test generation.: The mutation-testing literature is cited in detail in [1]. The relationship is: that work measured what specifications do for LLM-generated tests, and the present work makes the specifications transportable to LLMs and other consumers that do not have the source.

IX. CONCLUSION

We asked whether automatically inferred specifications can be embedded in compiled Java artefacts and REST API descriptions with sufficient fidelity that downstream consumers can use them as if they had the source. We answered in the affirmative, with two specific constructions:

- 1) A runtime-retained Java annotation, `@JmlSpecs`, written into class-file `RuntimeVisibleAnnotations` attributes by an ASM-based embedder. Roundtrip fidelity is 100.0% on three real-world libraries (Apache Commons Lang, Apache Commons IO, Guava — totalling 20,546 methods) and on a real-inferred-specs corpus of 10,332 specs across the same three libraries’ source trees, spanning quantifiers, `\old` references, and branch-conditional implications. Bytecode-size overhead after the format-optimisation step described in Section V-F is 5.27–6.97% on synthetic specs and 4.63–10.20% on real inferred specs (a 48–55% reduction over the initial per-clause-annotation format on the synthetic corpora, 31% on the only real-inference corpus measured under v1). Embedding throughput is 6,000–22,000 methods/second; reading throughput 74,000–328,000 methods/second.
- 2) An OpenAPI 3.x extension family (`x-jml-requires`, `x-jml-ensures`, `x-jml-assignable`, `x-jml-signals`, `x-jml-invariant`) per-operation and per-schema, plus a sidecar JSON fallback for documents generated by frameworks. The extension is preserved by standard OpenAPI parsers without modification.

Our negative test confirms that classes embedded with `@JmlSpecs` load correctly in JVMs that have no awareness of the annotation type — the embedder is consumer-tolerant by design. Together with the inference engine of [1], the path from source code to consumer-accessible contract is complete: the inferer produces specifications, the embedder ships them alongside compiled artefacts, and the OpenAPI extension carries them across service boundaries.

We see three threads of follow-up work. First, validation across the remaining seven libraries identified in our protocol (Apache Commons Math, jOOL, Vavr, jsoup, JFreeChart, JUnit, AssertJ). Second, integration with the LLM-based test generator of [1], which can now consume specifications directly from the JAR rather than requiring source. Third, an extension of the inference itself to compose specifications across method boundaries; the carrier built here is what makes that composition transportable to consumers.

The replication package, including the full source of the embedder/extractor pair, the OpenAPI extension JSON Schema, and the empirical validation harness, accompanies this paper.

DATA AVAILABILITY

A replication package containing the embedder/extractor source, the OpenAPI extension JSON Schema, raw measurement logs, and reproduction scripts is available at [https://github.com/\[anonymized\]/jml-inferer](https://github.com/[anonymized]/jml-inferer).

REFERENCES

- [1] B. Edmonds and M. Utting, “Do specifications help LLMs write better unit tests? an empirical study of specification-guided test generation with heuristically inferred JML,” 2026, manuscript under review at the

- Journal of Software: Evolution and Process. Replication package: <https://github.com/ziggyfish/jml-inferer>.
- [2] B. Meyer, “Applying “design by contract”,” *Computer*, vol. 25, no. 10, pp. 40–51, 1992.
- [3] G. T. Leavens, A. L. Baker, and C. Ruby, “Preliminary design of JML: A behavioral interface specification language for Java,” *ACM SIGSOFT Software Engineering Notes*, vol. 31, no. 3, pp. 1–38, 2006.
- [4] G. T. Leavens *et al.*, “JML reference manual,” <http://www.jmlspecs.org>, 2013, accessed: 2025-01-15.
- [5] D. R. Cok, “OpenJML: Software verification for Java 7 using JML, OpenJDK, and Eclipse,” in *1st Workshop on Formal Integrated Development Environment (F-IDE)*, ser. EPTCS, vol. 149, 2014, pp. 79–92.
- [6] M. D. Ernst, J. H. Perkins, P. J. Guo, S. McCamant, C. Pacheco, M. S. Tschantz, and C. Xiao, “The Daikon system for dynamic detection of likely invariants,” *Science of Computer Programming*, vol. 69, no. 1–3, pp. 35–45, 2007.
- [7] M. Barnett and K. R. M. Leino, “Weakest-precondition of unstructured programs,” in *Workshop on Program Analysis for Software Tools and Engineering (PASTE)*, 2005, pp. 82–87.
- [8] A. Astorga, S. Srisakaokul, X. Xiao, and T. Xie, “PreInfer: Automatic inference of preconditions via symbolic analysis,” in *48th Annual IEEE/IFIP International Conference on Dependable Systems and Networks (DSN)*. IEEE, 2018, pp. 678–689.
- [9] A. Blasi, A. Goffi, K. Kuznetsov, A. Gorla, M. D. Ernst, M. Pezzè, and S. D. Castellanos, “Translating code comments to procedure specifications,” in *Proceedings of the 27th ACM SIGSOFT International Symposium on Software Testing and Analysis*, ser. ISSTA ’18. ACM, Jul. 2018, pp. 242–253. [Online]. Available: <http://dx.doi.org/10.1145/3213846.3213872>
- [10] OW2 Consortium, “ASM: a Java bytecode manipulation and analysis framework,” <https://asm.ow2.io/>, 2024, accessed: 2026-05-19.
- [11] Pact Foundation, “Pact: consumer-driven contract testing,” <https://docs.pact.io/>, 2024, accessed: 2026-05-19.
- [12] A. Lercher, C. Macho, C. Bauer, and M. Pinzger, “Generating accurate OpenAPI descriptions from Java source code,” in *Proceedings of the 32nd IEEE International Conference on Software Analysis, Evolution and Reengineering (SANER)*, 2025, arXiv preprint: <https://arxiv.org/abs/2410.23873>.